

基于K8S的 iLogtail/Filebeat+KAFKA+Vector+ClickHouse/ES大规模日志采集最佳实践

本付费合集至少包括以下内容

- 基于Vector + ClickHouse的请求日志采集与Grafana日志分析看板实战
- 基于ClickHouse的请求日志，URI级别的5XX与耗时异常监控告警实践
- 基于K8S的iLogtail/Filebeat + KAFKA + Vector + ClickHouse/ES大规模日志采集方案
 - 基于K8S的KRaft模式KAFKA集群部署方案
 - 基于K8S的ClickHouse集群部署方案
 - 基于K8S的ES集群部署方案
- **VIP微信群：提供运维交流答疑，原创运维技术文章优先本集合内发布，合集内文章疑难解答与远程协助。**

为什么要创建付费合集

- 付费合集中包含的许多内容是独家提供的，只有付费用户才能享受到这些独特的资源和信息。
- 付费合集汇集了系统化的专题内容，读者可以更全面地掌握某一领域的知识，实现更深层次的学习和理解。
- 付费不仅是对内容的认可，也是对创作者辛勤付出的支持，帮助他们持续创作更多优质内容！

目录

基于K8S的iLogtail/Filebeat+KAFKA+Vector+ClickHouse/ES大规模日志采集最佳实践

本付费合集至少包括以下内容

为什么要创建付费合集

目录

整体架构

采集源

采集组件

iLogtail

Filebeat

Vector

消息组件：KAFKA

什么情况下需要使用KAFKA？

使用KAFKA的好处

日志聚合处理与入库

Vector

iLogtail

存储数据库

ElasticSearch

ClickHouse

展示分析

Kibana


ClickVisual

Grafana

💎【采集器】部署与日志的采集配置

- iLogtail: DaemonSet部署
- iLogtail: 采集配置
 - 采集容器标准输出并写入KAFKA
 - 采集容器内文件日志并写入KAFKA
- Filebeat: DaemonSet部署
- Filebeat: 采集配置
 - 采集容器文件日志和标准输出并写入KAFKA
- ◆【消息组件】KAFKA部署与配置
- ◆【处理器】部署与日志处理配置
 - Vector: Deployment部署
 - Vector: 日志处理与入库配置
 - 消费KAFKA处理日志并入库配置
 - iLogtail: 日志处理与入库配置
 - 采集容器标准输出，处理JAVA多行日志并直接写入ES8
- ◆【数据库】ClickHouse与ES部署

整体架构



采集源

- K8S容器内文件日志
- K8S容器标准输出

采集组件

iLogtail

- 部署方式Daemonset
- 支持采集任意K8S容器内文件日志，提取K8S元信息
- 支持采集容器标准输出
- 不支持写入ES7.X
- 占用CPU比Filebeat大幅降低，占用内存与Filebeat相当

Filebeat

- 部署方式Daemonset
- 支持采集K8S容器内在 `emptyDir` 目录下的文件日志，提取K8S元信息
- 支持采集容器标准输出

Vector

- 部署方式Daemonset
- 支持采集容器标准输出

消息组件：KAFKA

什么情况下需要使用KAFKA？

- 日志采集组件、日志处理组件出现使用CPU、内存异常升高
- 日志入库的延迟过长。
- 出现大量入库失败的情况。

使用KAFKA的好处

- 高吞吐量与低延迟：支持高并发写入和读取，实现实时数据处理和分析。
- 暂存机制：在数据入库失败时，可以配置系统从KAFKA重新消费这些数据。
- 采集和入库解耦：生产者只需将日志发送到KAFKA，不需要关心谁来消费或消费的速度。消费者可以独立地从KAFKA读取数据，处理和扩展变得更加灵活。
- 削峰填谷：即使在日志生成高峰期，KAFKA也能暂时存储大量数据，防止下游系统因处理能力不足而崩溃。消费者可以以自己的节奏处理数据，KAFKA通过持久化存储和批量传输来平滑数据流，避免系统在高负载时过载。

日志聚合处理与入库

Vector

- Rust
- 支持IP地址库解析，UserAgent解析
- 强大的数据处理函数、智能均衡KAFKA分区消费与入库，官测正则处理13.2M/s
- 比logstash强悍太多的性能，声称比同类方案快10 倍

iLogtail

- C++ & Golang
- 不支持IP地址库解析，不支持UserAgent解析
- 原生插件支持C++加速，宣称正则处理20M/s（原生插件仅可用于采集文本日志），其它类型可使用扩展插件处理，但性能会有一定影响。

存储数据库

ElasticSearch

全文检索引擎的文档数据库，对于**业务日志、异常日志、多行日志**这类，**非结构化**、半结构化的日志数据，经常需要做关键字查询，**模糊匹配**等操作，非常适合使用ES，使用倒排索引实现快速**全文搜索**。存储压缩率低，CPU/内存占用高。聚合分析性能没有ClickHouse好。

ClickHouse

一个列式存储数据库，尤其擅长处理**结构化**的大规模的SQL查询和聚合分析操作，所以针对NGINX这类结构化的**请求日志**，在处理**多维分析、聚合查询、分组统计**等操作速度极快，并且压缩比极高，存储成本比ES低10倍，CPU、内存的占用也有巨大优势。全文搜索性能没有ES好。

展示分析

Kibana

- 支持ES，日志查询分析方便易用
- 查询页的分组统计默认只针对前500行（不能针对所有行）
- 可查询指定行的上下文

ClickVisual

- 支持ClickHouse，类似Kibana界面，没有Kibana好用
- 查询页的分组统计针对所有行
- 不可查询指定行的上下文

Grafana

- 支持ClickHouse与ES，数据分析展示最佳，日志查询界面没有Kibana好用。

💎【采集器】部署与日志的采集配置

iLogtail: DaemonSet部署

```
1  apiVersion: apps/v1
2  kind: DaemonSet
3  metadata:
4    name: ilogtail-ds
5    namespace: ops-monit
6    labels:
7      k8s-app: logtail-ds
8  spec:
9    selector:
10     matchLabels:
11       k8s-app: logtail-ds
12    template:
13     metadata:
14       labels:
15         k8s-app: logtail-ds
16     spec:
17       tolerations:
18         - operator: Exists
19       containers:
20         - name: logtail
21           command:
22             - /usr/local/ilogtail/ilogtail_control.sh
23           args:
24             - "start_and_block"
25             - "-enable_containerd_upper_dir_detect=true"
26             - "-dirfile_check_interval_ms=5000"
```

```

27     - "-logtail_checkpoint_check_gc_interval_sec=120"
28     env:
29       - name: _node_name_
30         valueFrom:
31           fieldRef:
32             apiVersion: v1
33             fieldPath: spec.nodeName
34       - name: _node_ip_
35         valueFrom:
36           fieldRef:
37             apiVersion: v1
38             fieldPath: status.hostIP
39       - name: cpu_usage_limit
40         value: "1"
41       - name: mem_usage_limit
42         value: "512"
43     image: >-
44     sls-opensource-registry.cn-shanghai.cr.aliyuncs.com/ilogtail-
community-edition/ilogtail:latest
45     imagePullPolicy: IfNotPresent
46     resources:
47       limits:
48         cpu: 1000m
49         memory: 1Gi
50       requests:
51         cpu: 40m
52         memory: 38Mi
53     volumeMounts:
54       - mountPath: /var/run
55         name: run
56       - mountPath: /logtail_host
57         mountPropagation: HostToContainer
58         name: root
59         readOnly: true
60       - mountPath: /usr/local/ilogtail/checkpoint
61         name: checkpoint
62       - mountPath: /usr/local/ilogtail/config/local
63         name: user-config
64         readOnly: true
65     dnsPolicy: ClusterFirstWithHostNet
66     hostNetwork: true
67     volumes:
68       - hostPath:
69         path: /var/run
70         type: Directory
71         name: run
72       - hostPath:
73         path: /
74         type: Directory
75         name: root
76       - hostPath:
77         path: /etc/ilogtail-ilogtail-ds/checkpoint
78         type: DirectoryOrCreate
79         name: checkpoint
80       - configMap:
81         defaultMode: 420

```

```
82     name: ilogtail-user-cm
83     name: user-config
```

iLogtail: 采集配置

采集容器标准输出并写入KAFKA

根据微服务命名空间与标签来匹配微服务

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: ilogtail-user-cm
5    namespace: ops-monit
6  data:
7    nginx_stdout.yaml: |
8      enable: true
9      inputs:
10       - Type: service_docker_stdout # 采集K8S容器标准输出
11         Stdout: true
12         Stderr: false
13         K8sNamespaceRegex: "^istio-system$" # 指定命名空间
14         IncludeK8sLabel:
15           app: istio-ingressgateway # 指定需要匹配的标签
16       flushers:
17       - Type: flusher_kafka_v2
18         Brokers:
19           - ops-log1.kafka.casstime.com:9092
20           - ops-log2.kafka.casstime.com:9092
21           - ops-log3.kafka.casstime.com:9092
22         Topic: prod-istio-std
23         MaxMessageBytes: 10000000
```

采集容器内文件日志并写入KAFKA

根据文件在容器内的路径来匹配

```
1  #接上面的配置，支持在一个configmap中有多个配置文件。
2  java_file.yaml: |
3    enable: true
4    inputs:
5      - Type: input_file
6        FilePaths:
7          - /opt/logs/*.log # 指定容器内的日志路径
8        EnableContainerDiscovery: true # 表示采集容器内的文件日志
9        ExcludeFiles:
10         - access.* # 需要排除的文件名
11        ContainerFilters:
12          ExcludeK8sLabel:
13            app: pod-agentbuy-service # 可排除包含某些标签的容器
14        Multiline:
15          StartPattern: \d+-\d+-\d+.* # 匹配多行日志开头的正则
16          AppendingLogPositionMeta: true # 添加采集文件的元信息
17        processors:
18          - Type: processor_add_fields # 增加字段
```

```

19     Fields:
20         logtype: java-file
21         env: cassmall
22     flushers:
23         - Type: flusher_kafka_v2
24         Brokers:
25             - cassops-log1.kafka.casstime.com:9092
26             - cassops-log2.kafka.casstime.com:9092
27             - cassops-log3.kafka.casstime.com:9092
28         Topic: prod-java-file-ilogtail
29         MaxMessageBytes: 10000000
30         Convert:
31             TagFieldsRename: # 对K8S元数据的字段名重命名
32             k8s.pod.name: 'pod_name'
33             k8s.node.name: 'node_name'
34             k8s.namespace.name: 'namespace'
35             container.image.name: 'deployimage'
36             log.file.path: 'path'

```

Filebeat: DaemonSet部署

```

1  apiVersion: v1
2  kind: ServiceAccount
3  metadata:
4      name: filebeat
5      namespace: ops-monit
6  ---
7  apiVersion: rbac.authorization.k8s.io/v1
8  kind: ClusterRole
9  metadata:
10     annotations:
11     labels:
12         k8s-app: filebeat
13     name: filebeat
14  rules:
15  - apiGroups:
16      - ""
17    resources:
18      - namespaces
19      - pods
20    verbs:
21      - get
22      - watch
23      - list
24  apiVersion: rbac.authorization.k8s.io/v1
25  ---
26  kind: ClusterRoleBinding
27  metadata:
28     annotations:
29     name: filebeat
30  roleRef:
31     apiGroup: rbac.authorization.k8s.io
32     kind: ClusterRole
33     name: filebeat
34  subjects:

```

```
35 - kind: ServiceAccount
36   name: filebeat
37   namespace: ops-monit
38 ---
39 apiVersion: apps/v1
40 kind: DaemonSet
41 metadata:
42   labels:
43     k8s-app: filebeat
44   name: filebeat
45   namespace: ops-monit
46 spec:
47   selector:
48     matchLabels:
49       k8s-app: filebeat
50   template:
51     metadata:
52       labels:
53         k8s-app: filebeat
54     spec:
55       containers:
56       - args:
57         - -c
58         - /etc/filebeat.yml
59         - -e
60       env:
61       - name: NODE_NAME
62         valueFrom:
63           fieldRef:
64             apiVersion: v1
65             fieldPath: spec.nodeName
66       image: elastic/filebeat:7.17.24
67       #image: elastic/filebeat:7.6.2
68       imagePullPolicy: IfNotPresent
69       name: filebeat
70       resources:
71         limits:
72           cpu: "1"
73           memory: 600Mi
74         requests:
75           cpu: 100m
76           memory: 120Mi
77       securityContext:
78         runAsUser: 0
79       terminationMessagePath: /dev/termination-log
80       terminationMessagePolicy: File
81       volumeMounts:
82       - mountPath: /etc/filebeat.yml
83         name: config
84         readOnly: true
85         subPath: filebeat.yml
86       - mountPath: /usr/share/filebeat/data
87         name: data
88       - mountPath: /var/lib/docker/containers
89         name: varlibdockercontainers
90         readOnly: true
```



```

91     - mountPath: /var/log
92       name: varlog
93       readOnly: true
94     - mountPath: /var/lib/kubelet/pods
95       name: pod-test
96       readOnly: true
97   dnsConfig:
98     options:
99       - name: single-request-reopen
100   dnsPolicy: ClusterFirstWithHostNet
101   restartPolicy: Always
102   securityContext: {}
103   serviceAccount: filebeat
104   serviceAccountName: filebeat
105   tolerations:
106     - operator: Exists
107   volumes:
108     - configMap:
109         defaultMode: 416
110         name: filebeat-config
111         name: config
112     - hostPath:
113         path: /var/lib/docker/containers
114         type: ""
115         name: varlibdockercontainers
116     - hostPath:
117         path: /var/log
118         type: ""
119         name: varlog
120     - hostPath:
121         path: /var/lib/filebeat-data
122         type: DirectoryOrCreate
123         name: data
124     - hostPath:
125         path: /mnt/paas/kubernetes/kubelet/pods
126         type: DirectoryOrCreate
127         name: pod-test

```

Filebeat: 采集配置

采集容器文件日志和标准输出并写入KAFKA

注意容器内的文件日志必须存放在 `emptyDir` 目录下

```

1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: filebeat-config
5    namespace: ops-monit
6    labels:
7      k8s-app: filebeat
8  data:
9    filebeat.yml: |-
10     filebeat.inputs:
11     - type: log # 文件日志

```

```

12     paths:
13         - /var/lib/kubelet/pods/*/volumes/kubernetes.io~empty-dir/**/*.*log #
默认empty-dir在节点主机的路径
14     exclude_files: ['.*access.*'] # 可以排除某些文件名
15     multiline.type: pattern # 多行日志匹配
16     multiline.pattern: ^\d{4}-\d{2}-\d{2}\s\d{2}:\d{2}:\d{2}
17     multiline.negate: true
18     multiline.match: after
19     fields: # 增加字段
20         logtype: java-file
21         env: f2b-gamma
22     processors:
23         - add_kubernetes_metadata: # 增加K8S元数据
24             indexers:
25                 - pod_uid:
26             matchers:
27                 - logs_path:
28                     logs_path: '/var/lib/kubelet/pods/'
29                     resource_type: 'pod'
30
31     - type: container # 采集容器标准输出
32     paths:
33         - /var/log/containers/*_f2b_*node*.log # 默认标准输出在节点主机的路径
34     fields: # 增加字段
35         logtype: json-std
36         env: f2b-gamma
37     processors:
38         - add_kubernetes_metadata: # 增加K8S元数据
39             matchers:
40                 - logs_path:
41                     logs_path: /var/log/containers
42
43     output.kafka: # 输出到KAFKA
44         hosts: ["10.1.230.152:9092", "10.1.227.35:9092", "10.1.247.50:9092"]
45         topic: 'topic-gamma-logs'
46         max_message_bytes: 10000000
47         partition.round_robin:
48             reachable_only: false

```

💎 【消息组件】KAFKA部署与配置

💎 【处理器】部署与日志处理配置

Vector: Deployment部署

```

1  apiVersion: v1
2  kind: ServiceAccount
3  metadata:
4      namespace: ops-monit
5      name: vector
6  automountServiceAccountToken: true
7  ---
8  apiVersion: rbac.authorization.k8s.io/v1

```

```

9 kind: ClusterRole
10 metadata:
11   name: vector
12   labels:
13     app.kubernetes.io/name: vector
14 rules:
15   - apiGroups:
16     - ""
17     resources:
18     - namespaces
19     - nodes
20     - pods
21     verbs:
22     - list
23     - watch
24 ---
25 apiVersion: rbac.authorization.k8s.io/v1
26 kind: ClusterRoleBinding
27 metadata:
28   name: vector
29   labels:
30     app.kubernetes.io/name: vector
31 roleRef:
32   apiGroup: rbac.authorization.k8s.io
33   kind: ClusterRole
34   name: vector
35 subjects:
36   - kind: ServiceAccount
37     name: vector
38     namespace: ops-monit
39 ---
40 apiVersion: apps/v1
41 kind: Deployment
42 metadata:
43   name: vector
44   namespace: ops-monit
45   labels:
46     app.kubernetes.io/name: vector
47 spec:
48   replicas: 1
49   selector:
50     matchLabels:
51       app.kubernetes.io/name: vector
52   template:
53     metadata:
54       annotations: {}
55       labels:
56         app.kubernetes.io/name: vector
57         vector.dev/exclude: "true"
58     spec:
59       serviceAccountName: vector
60       dnsPolicy: ClusterFirst
61       containers:
62       - name: vector
63         image: "registry.cn-shenzhen.aliyuncs.com/stars1/vector:0.42.0-
alpine"

```

```
64     imagePullPolicy: IfNotPresent
65     args:
66       - --config-dir
67       - /etc/vector/
68     env:
69       - name: TZ
70         value: Asia/Shanghai
71     ports:
72       - name: api
73         containerPort: 8686
74         protocol: TCP
75     resources:
76       limits:
77         cpu: '2'
78         memory: 2Gi
79       requests:
80         cpu: 500m
81         memory: 100Mi
82     volumeMounts:
83       - name: data
84         mountPath: "/vector-data-dir"
85       - name: config
86         mountPath: "/etc/vector/"
87         readOnly: true
88     terminationGracePeriodSeconds: 60
89     affinity:
90       podAntiAffinity:
91         requiredDuringSchedulingIgnoredDuringExecution:
92           - labelSelector:
93               matchLabels:
94                 app.kubernetes.io/name: vector
95             topologyKey: kubernetes.io/hostname
96     volumes:
97       - name: data
98         emptyDir: {}
99       - name: config
100         projected:
101           sources:
102             - configMap:
103                 name: vector
104     ---
105     apiVersion: v1
106     kind: Service
107     metadata:
108       namespace: ops-monit
109       name: vector
110       labels:
111         app.kubernetes.io/name: vector
112     spec:
113       ports:
114         - name: api
115           port: 8686
116           protocol: TCP
117       selector:
118         app.kubernetes.io/name: vector
119       type: NodePort
```

Vector: 日志处理与入库配置

消费KAFKA处理日志并入库配置

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: vector
5    namespace: ops-monit
6  data:
7    global.toml: |
8      data_dir = "/vector-data-dir"
9      timezone = "Asia/Shanghai"
10     [api]
11       address = "0.0.0.0:8686"
12       enabled = true
13
14     prod-java-file-ilogtail.toml: |- # java业务日志入库ES(toml格式配置文件写法)
15       [sources.1kafka_java_ilogtail]
16         type = "kafka"
17         bootstrap_servers = "10.0.100.1:9092,10.0.100.2:9092,10.0.100.3:9092"
18         group_id = "prod-java-file-ilogtail"
19         topics = [ "prod-java-file-ilogtail" ]
20         decoding.codec = "json"
21         commit_interval_ms = 1000
22
23       [transforms.2parse_java_all_ilogtail]
24       inputs = ["1kafka_java_ilogtail"]
25       type = "remap"
26       source = '''
27         .namespace = .tags.namespace
28         .podname = .tags.pod_name
29         .nodename = .tags.node_name
30         .appname = .tags.appname
31         .deployimage = .tags.deployimage
32         # java日志使用正则解析,获取时间和日志级别字段即可,并保留原始日志.
33         . |= parse_regex!(.fields.content, r'^(?P<time>\d+-\d+-\d+
\d+:\d+:\d+,\d+)(\s|:)+(?P<level>\s+)(\s|\s)*$')
34         '''
35
36       [transforms.3delfields_java_ilogtail]
37       inputs = ["2parse_java_all_ilogtail"]
38       type = "remap"
39       source = '''
40         .unix_time, err = to_int(format_timestamp!(.timestamp, format: "%s"))
41         .utc8_time = .unix_time + 28800
42         .index_day = format_timestamp!(to_timestamp!(.utc8_time), "%Y.%m.%d")
43         .path = .tags.path
44         del(.unix_time)
45         del(.utc8_time)
46         del(.tags)
47         .message = del(.fields.content)
48         del(.topic)
49         del(.offset)
```

```

50     del(.partition)
51     del(.message_key)
52     del(.source_type)
53     del(.fields.__file_offset__)
54     .@timestamp = parse_timestamp(.time, format: "%F %T,%3f") ?? now()
55     .capturetime = del(.timestamp)
56     del(.time)
57     ns, err = string(.namespace)
58     if err != null {
59         .namespace = "null"
60     }
61     '''
62
63     [sinks.9es_java_ilogtail]
64         type = "elasticsearch"
65         inputs = ["3delfields_java_ilogtail"]
66         endpoint = "http://10.1.72.33:19200"
67         auth.strategy = "basic"
68         auth.user = "elastic"
69         auth.password = "M1k2w12Ig4A"
70         mode = "bulk"
71         buffer.max_events = 50000
72         batch.max_events = 50000
73         batch.timeout_secs = 5
74         bulk.index = "hwprod-java-file-{{ namespace }}-{{ index_day }}"
75
76     istio.yaml: |- # istio请求日志入库ClickHouse(yaml格式配置文件写法)
77     sources:
78         01_kafka_hwprod_istio:
79             type: "kafka"
80             bootstrap_servers: "kafka1:9092,kafka2:9092,kafka3:9092"
81             group_id: "hwprod_istio"
82             topics: [ "prod-istio-std" ]
83             decoding:
84                 codec: "json"
85             commit_interval_ms: 1000
86
87     transforms:
88         02_parse_hwprod_istio:
89             drop_on_error: true
90             reroute_dropped: true
91             type: remap
92             inputs:
93                 - 01_kafka_hwprod_istio
94             source: |
95                 # json日志直接解析json,即可获取所有字段.
96                 . = parse_json!(.contents.content)
97                 .start_time = to_unix_timestamp(parse_timestamp!(.start_time,
format: "%+"),unit: "milliseconds")
98                 .upstream_service_time = to_int!(.upstream_service_time)
99
100     sinks:
101         03_ck_hwprod_istio:
102             type: clickhouse
103             inputs:
104                 - 02_parse_hwprod_istio

```

```

105     database: istio logs
106     endpoint: http://10.7.0.226:28123
107     table: hwprod_istio_local
108     compression: gzip
109 04_out_istio_dropped:
110     type: blackhole
111     inputs:
112         - 02_parse_hwprod_istio.dropped

```

iLogtail：日志处理与入库配置

采集容器标准输出，处理JAVA多行日志并直接写入ES8

```

1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: ilogtail-user-cm
5    namespace: ops-monit
6  data:
7    java_std_es8.yaml: |
8      enable: true
9      inputs:
10         - Type: service_docker_stdout
11           Stderr: true
12           Stdout: true
13           BeginLineCheckLength: 10
14           BeginLineRegex: "\\d+-\\d+-\\d+.*" # JAVA多行日志开头匹配正则
15           K8sNamespaceRegex: "^ (saas-prod) $"
16     processors:
17         - Type: processor_string_replace # 如果日志的标准输出有颜色,可以使用此正则替
      换掉
18           SourceKey: content
19           Method: regex
20           Match: "\\d*\\;.*\\d*m|\\u001b"
21           ReplaceString: ''
22         - Type: processor_regex # JAVA日志处理: 只从日志中提取时间和级别字段, 并保留
      源。
23           SourceKey: content
24           Regex: "(\\d{4}-\\d{2}-\\d{2} \\d{2}:\\d{2}:\\d{2},\\d{3}) .*[0-9a-z]*([A-
      z]+)"
25           Keys:
26             - log_time
27             - level
28           KeepSource: true
29         - Type: processor_gotime # 日志中的时间字段转成ES中@timestamp标准日期时间格
      式。
30           SourceKey: "log_time"
31           SourceFormat: "2006-01-02 15:04:05,000"
32           SourceLocation: 8
33           DestKey: "@timestamp"
34           DestFormat: "2006-01-02T15:04:05.000Z"
35           DestLocation: 0
36           SetTime: false
37           KeepSource: false
38         - Type: processor_rename # 字段重命名

```

```
39     SourceKeys:
40         - _time_
41         - _source_
42     DestKeys:
43         - cap_time
44         - log_src
45     flushers:
46         #- Type: flusher_stdout
47         # OnlyStdout: true
48         - Type: flusher_elasticsearch
49     Addresses:
50         - http://10.26.6.152:9200
51         - http://10.26.6.153:9200
52         - http://10.26.6.154:9200
53     Convert:
54         Protocol: custom_single_flatten # 数据字段做一级打平
55         Encoding: json
56     Index: java-logs-%{tag.k8s.namespace.name}_%{+yyyy.MM.dd}
57     Authentication:
58         PlainText:
59             Username: ilogtail1
60             Password: fSdgDh4jfg
61     HTTPConfig:
62         MaxIdleConnsPerHost: 10
63         ResponseHeaderTimeout: Second
```

【数据库】ClickHouse与ES部署